Lab 10 — Image Segmentation and Feature Extraction

Course: ARTI404 – Image Processing
College of Computer Science and Information Technology
Imam Abdulrahman Bin Faisal University

This notebook contains:

1. The two examples from the lab manual (Canny Edge Detection, Harris Corner Detection).
2. The assessment task (Otsu's thresholding).

Name: Abdulsamad Mousa Tuhaishi

ID: 2240001093

Setup — Import libraries

In [1]:
```
import cv2
import numpy as np
import matplotlib.pyplot as plt
from skimage import data   # provides sample images (camera, checkerboard

# Make plots show inside the notebook
%matplotlib inline
```

Task 1 — Canny Edge Detection

An **edge** is where brightness changes sharply. Canny (1986) finds edges in 5 steps: Gaussian blur → gradients → non-maximum suppression (thin the edges) → double thresholding → hysteresis (keep weak edges only if they connect to strong ones).

`cv2.Canny(img, T_low, T_high)`:

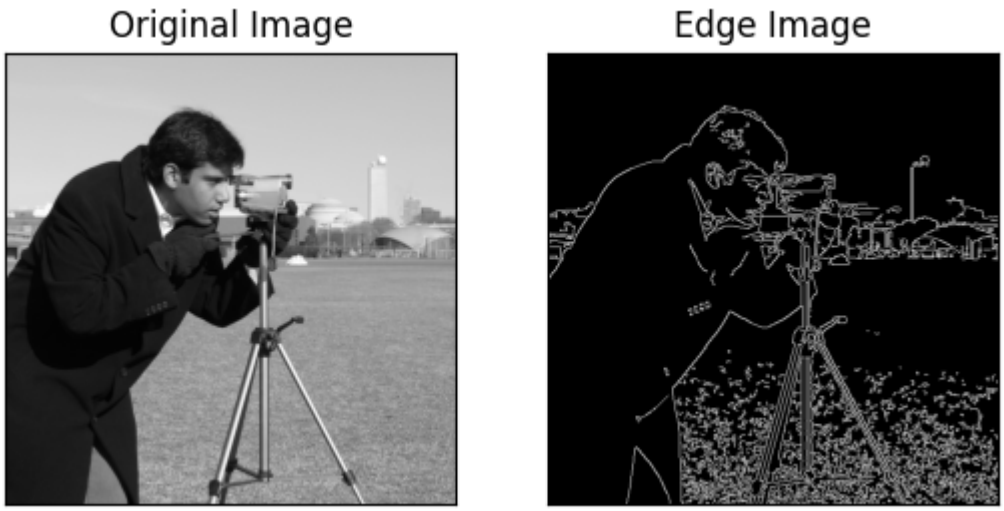
- `T_low` = lower threshold (weak edges)
- `T_high` = upper threshold (strong edges)
- Rule of thumb: `T_high ≈ 2×T_low` or `3×T_low`.

In [2]:
```
# --- Code from the lab manual ---
img = data.camera()

edges = cv2.Canny(img, 100, 200)

plt.subplot(121), plt.imshow(img, cmap='gray')
plt.title('Original Image'), plt.xticks([]), plt.yticks([])
plt.subplot(122), plt.imshow(edges, cmap='gray')
```

```
plt.title('Edge Image'), plt.xticks([]), plt.yticks([])
plt.show()
```



Task 2 — Harris Corner Detection

A **corner** is a point where the image changes in *two* different directions — which makes corners highly distinctive (useful for matching, tracking, camera calibration).

- Large positive `R` → corner
- Large negative `R` → edge
- Small `R` → flat region

Parameters: `blockSize` (window size), `ksize` (Sobel kernel size), `k` (sensitivity, typically 0.04–0.06).

> **Note:** OpenCV uses BGR order, so `(0, 0, 255)` is red.

In [3]:
```
# --- Code from the lab manual ---
board = data.checkerboard()
image = cv2.cvtColor(board, cv2.COLOR_GRAY2BGR)
image = cv2.resize(image, (400, 400), interpolation=cv2.INTER_NEAREST)

# Convert to grayscale
gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
gray = np.float32(gray)

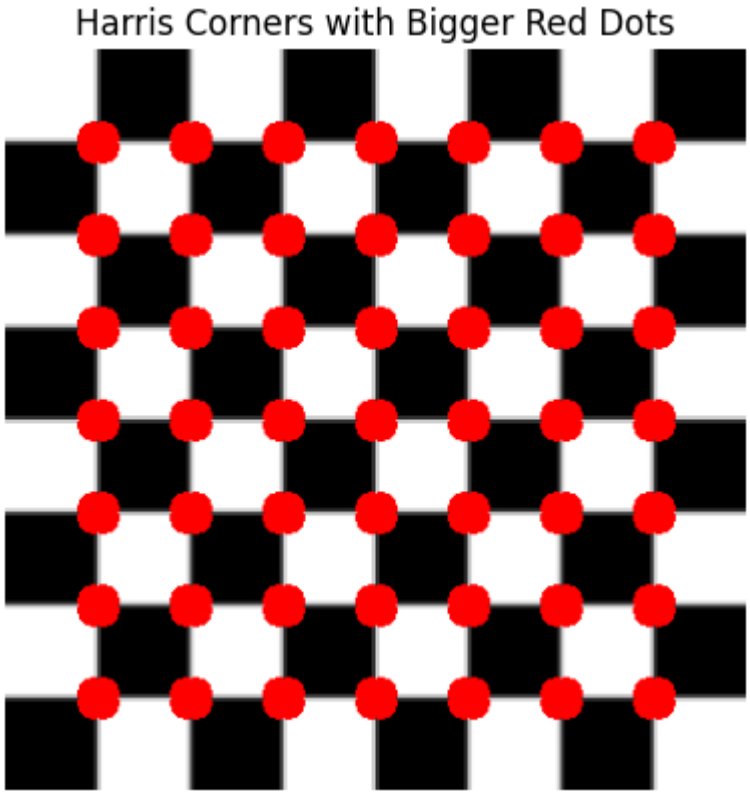
# Harris Corner Detection
harris = cv2.cornerHarris(gray, blockSize=2, ksize=3, k=0.04)
harris = cv2.dilate(harris, None)   # to mark corners more clearly

# Threshold for detecting strong corners
threshold = 0.01 * harris.max()
corner_coords = np.argwhere(harris > threshold)

# Draw larger red dots (radius = 7)
for y, x in corner_coords:
    cv2.circle(image, (x, y), radius=7, color=(0, 0, 255), thickness=-1)

# Display
plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB))
```

```
plt.title('Harris Corners with Bigger Red Dots')
plt.axis('off')
plt.show()
```



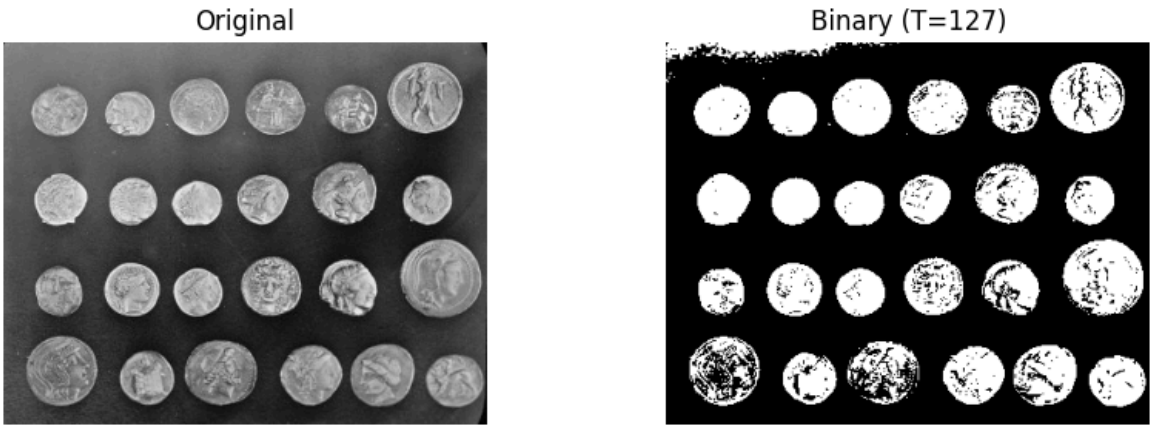
Simple (Global) Binary Thresholding

The most basic segmentation: pick a fixed threshold `T` and turn every pixel above it white, below it black.

In [4]:
```
img = data.coins()

T = 127
_, binary = cv2.threshold(img, T, 255, cv2.THRESH_BINARY)

plt.figure(figsize=(9, 3))
plt.subplot(1, 2, 1), plt.imshow(img, cmap='gray'), plt.title('Original')
plt.subplot(1, 2, 2), plt.imshow(binary, cmap='gray'), plt.title(f'Binary
plt.tight_layout(), plt.show()
```



Out[4]: (None, None)

Q1 — Otsu's Thresholding (Assessment Task)

Otsu picks `T` **automatically** by maximizing the between-class variance — assumes the histogram has two peaks (foreground + background).

show the original image, its histogram, and the thresholded image.

In [5]:
```
img = data.coins()

T_otsu, otsu = cv2.threshold(img, 0, 255, cv2.THRESH_BINARY + cv2.THRESH_
print('Otsu threshold =', T_otsu)

plt.figure(figsize=(12, 3.5))

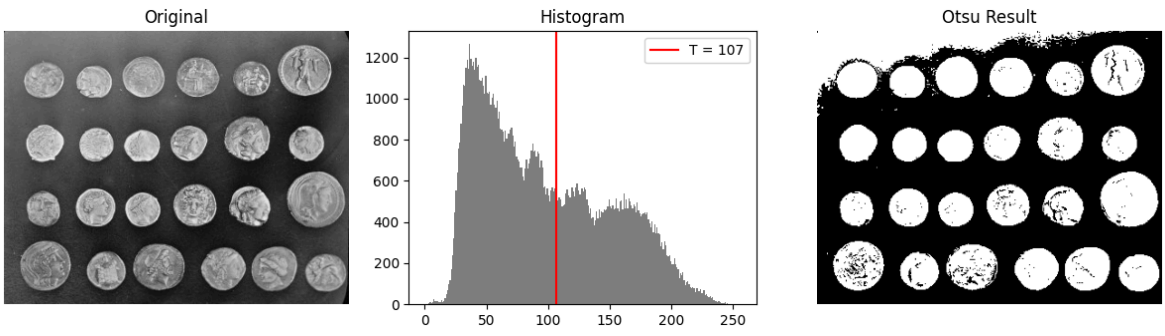
plt.subplot(1, 3, 1)
plt.imshow(img, cmap='gray')
plt.title('Original')
plt.axis('off')

plt.subplot(1, 3, 2)
plt.hist(img.ravel(), bins=256, range=(0, 256), color='gray')
plt.axvline(T_otsu, color='red', label=f'T = {T_otsu:.0f}')
plt.title('Histogram')
plt.legend()

plt.subplot(1, 3, 3)
plt.imshow(otsu, cmap='gray')
plt.title('Otsu Result')
plt.axis('off')

plt.tight_layout()
plt.show()
```

```
Otsu threshold = 107.0
```



Summary table

Technique	Category	What it finds	OpenCV function
Canny	Edge detection	Thin connected edges	`cv2.Canny`

Technique	Category	What it finds	OpenCV function
Sobel	Edge detection	Gradient magnitude (thick)	`cv2.Sobel`
Laplacian	Edge detection	Zero-crossings (2nd deriv)	`cv2.Laplacian`
Harris	Feature extraction	Corners	`cv2.cornerHarris`
Shi-Tomasi	Feature extraction	Corners (improved)	`cv2.goodFeaturesToTrack`
Simple threshold	Segmentation	Fixed-`T` binary mask	`cv2.threshold`
Otsu	Segmentation	Auto bimodal threshold	`cv2.threshold` + `THRESH_OTSU`
Adaptive threshold	Segmentation	Local-`T` binary mask	`cv2.adaptiveThreshold`

End of notebook.